Emanuele La Malfa

personal page

Work Experience

Research Associate Jul. 2023 – Current

Benchmarking Large Language Models. *Dept. of Computer Science, University of Oxford*

- Principal Investigators: Prof. Michael Wooldridge, Nigel Shadbolt, and Anthony Cohn.
- I conduct research on Large Language Models, with a particular focus on benchmarking their reasoning and planning capabilities.

Research Assistant Oct. 2019 – Mar. 2021

Enabling rapid adoption of artificial intelligence through an anonymized data protocol and explainable models. *University of Oxford, UK*

- Principal Investigator: Prof. Marta Kwiatkowska.
- Collaboration with GenieAI and funded through the InnovateUK scheme.
- The collaboration led to a paper published at EMNLP 2020.

Education

University of Oxford Oct. 2019 – Nov. 2023

PhD in Computer Science, supervised by Prof. Marta Kwiatkowska. *Oxford, UK*

Polytechnic University of Milan Feb. 2015 – Sept. 2017

Master's Degree in Computer Science and Engineering. *Milan, Italy*

Polytechnic University of Milan Sept. 2011 – Sept. 2014

Bachelor's Degree in Computer Engineering. *Milan, Italy*

Selected Publications

– 2025 –

Large Language Models Miss the Multi-Agent Mark
NeurIPS 2026 (position track, acceptance rate 6%)
Emanuele La Malfa, Gabriele La Malfa, Samuele Marro, Jie M. Zhang, Elizabeth Black, Michael Luck, Philip Torr and Michael Wooldridge

Language Models are Implicitly Continuous
ICLR 2025 (main track)
Samuele Marro, Davide Evangelista, X. Angelo Huang, **Emanuele La Malfa**, Michele Lombardi, Michael Wooldridge

One Language, Many Gaps: Evaluating Dialect Fairness and Robustness of Large Language Models in Reasoning Tasks
ACL 2025 (main track)
Fangru Li, Shaoguang Mao, **Emanuele La Malfa**, Valentin Hofmann, Adrian de Wynter, Jing Yao, Si-Qing Chen, Michael Wooldridge, Furu Wei

When Claims Evolve: Evaluating and Enhancing the Robustness of Embedding Models Against Misinformation Edits
ACL 2025 (Findings)
Jabez Magomere, **Emanuele La Malfa**, Manuel Tonneau, Ashkan Kazemi, Scott Hale

Understanding the Logical Capabilities of Large Language Models via Out-of-Context Representation Learning
EMNLP 2025 (Findings)
Jonathan Shaki, **Emanuele La Malfa**, Michael Wooldridge and Sarit Kraus

– 2024 –

Language-Models-as-a-Service: Overview of a New Paradigm and it its Challenges
Journal of Artificial Intelligence Research (JAIR) - **oral presentation at AAAI 2025** - media coverage here and here
Emanuele La Malfa, Aleks Petrov, Frieder Simon, Christoph Weinhuber, Raza Nazar, Anthony Cohn, Nigel Shadbolt and Michael Wooldridge

Graph-enhanced Large Language Models in Asynchronous Plan Reasoning
ICML 2024 (main track)
Fangru Lin, **Emanuele La Malfa**, Valentin Hofmann, Elle Michelle Yang, Anthony Cohn and Janet Pierrehumbert

Deep Neural Networks via Complex Network Theory: a Perspective
IJCAI 2024 (main track)
Emanuele La Malfa, Gabriele La Malfa, Giuseppe Nicosia, Vito Latora

A Notion of Complexity for Theory of Mind via Discrete World Models
EMNLP 2024 (Findings)
X. Angelo Huang, **Emanuele La Malfa**, Samuele Marro, Andrea Asperti, Anthony Cohn and Michael Wooldridge

– 2023-2020 –

Language Models Tokenizers Introduce Unfairness Between Languages
NeurIPS 2023 (main track) - website
Aleksandar Petrov, **Emanuele La Malfa**, Philip Torr, Adel Bibi

***The King is Naked*: on the Notion of Robustness for Natural Language Processing**
AAAI 2022 (main track) – **oral presentation** –
Emanuele La Malfa, Marta Kwiatkowska

On Guaranteed Optimal Robust Explanations for NLP Models
IJCAI 2021 (main track)
Emanuele La Malfa, Rhiannon Michelmore, Agnieszka Zbrzeny, Nicola Paoletti, Marta Kwiatkowska

Assessing Robustness of Text Classification through Maximal Safe Radius Computation
EMNLP 2020 (Findings)
Emanuele La Malfa, Min Wu, Luca Laurenti, Benjie Wang, Anthony Hartshorn, Marta Kwiatkowska

Pre-prints/Work Under Review

Fixed Point Explainability
Under review
Emanuele La Malfa, Jon Vadillo, Marco Molinari and Michael Wooldridge

Code Simulation as a Proxy for High-order Tasks in Large Language Models
Under review
Emanuele La Malfa, Christoph Weinhuber, Orazio Torre, Fangru Lin, Angelo X. Huang, Samuele Marro, Anthony Cohn, Nigel Shadbolt and Michael Wooldridge

Jailbreaking Large Language Models in Infinitely Many Ways
Technical report
Oliver Goldstein, **Emanuele La Malfa**, Felix Drinkall, Samuele Marro, Michael Wooldridge

A Scalable Communication Protocol for Networks of Large Language Models
Under review
Samuele Marro, **Emanuele La Malfa**, Jesse Wright, Guohao Li, Nigel Shadbolt, Michael Wooldridge, Philip Torr

Grants & Fellowships

Artificial Intelligence Safety Fund 450,000 USD
*AI Agent Evaluation & Synthetic Content RFP. PIs: **Emanuele La Malfa** and Samuele Marro.*

Schmidt AI2050 Senior Fellowship 1M USD
Foundations of LLM-based Multi-Agent Systems. PI: Prof. Michael Wooldridge. I wrote part of the grant proposal.

Teaching Experience

Deep Learning in Healthcare 2024 (HT)
Practical sessions *University of Oxford, UK*
Machine Learning 2023 (MT)
Classes *University of Oxford, UK*
Probabilistic Model Checking 2023 (MT)
Practical sessions *University of Oxford, UK*
Ethical Computing in Practice 2023 (HT)
Practical sessions *University of Oxford, UK*
Deep Learning in Healthcare 2023 (HT)
Practical sessions *University of Oxford, UK*
Probabilistic Model Checking 2022 (MT)
Practical sessions *University of Oxford, UK*
Machine Learning 2021 (MT)
Classes *University of Oxford, UK*
Probabilistic Model Checking 2021 (MT)
Practical sessions *University of Oxford, UK*
Fundamentals of Computer Science 2016 (Oct.-Dec.)
Practical sessions *Polytechnic University of Milan, Italy*

Conferences and Workshops Organization

Benchmarking Large Language Models 28/11/2023
Workshop Organizer *The Alan Turing Institute, London, UK*
LOD2020, **LOD2021**, **LOD2022**, **LOD2023**, **LOD2025**
General Chair (2025), Conference Chair (others) *Lake District/Siena*

Invited Lectures, Talks, and Presentations

Code Simulation Challenges for Large Language Models 02/02/24
Group Talk *Bocconi, Italy*

On Robustness for Natural Language Processing 19/04/2023
Group Talk *ICREA, Barcelona*

On the Notion of Robustness for Natural Language Processing 17/01/2023
Departmental Talk *King's College University of London, UK*

Robustness for Natural Language Processing 22/04/2022
Lecture – Deep Fridays *University of Bologna, Italy*

Explainable AI 04/03/2022
Lecture – Advanced Artificial Intelligence Course *Royal Holloway University of London, UK*

Academic Reviewing and Volunteering

The Alan Turing Institute - Reviewer October - December 2023
Reviewers for the Turing Fellow Program - Panel "Fundamental Research in Data Science and AI"

Ukrainian Global University - Interviewer April-June 2023
I interviewed Ukrainian students who want to study in a partner university abroad.

The Kharkiv and Przemyśl Project August 2022
I spent a week in Przemyśl (Poland) as a volunteer to help refugees who arrived (returned) from (to) Ukraine.

Ukrainian Global University - Interviewer June-May 2022
I interviewed a dozen of prospective undergraduate Ukrainian students who want to study in a partner university abroad.

Eutanasia Legale - Volunteer July 2021
I have collected signatures for a referendum to decriminalize euthanasia. The overall campaign gathered 1.2 million valid signatures.

Tutoring and Mentoring

Williams-Exeter Exchange Programme 2023-2025
Tutoring Saad Waheed, Alisa Kanganis, and Simon Socolow (Williams-Exeter Programme exchange students in machine learning).

University of Oxford 2022
Tutoring Edward Kusel and Aleksandar Radoslavov for their part-B projects (undergraduate in Computer Science).

Lead the Future - Mentor 2022-current
Lead the Future helps Italian STEM talents find their path to brilliant careers. I am currently mentoring 8 students.

Mentoring and Student Supervision

Samuele Marro
I co-supervised his Master's thesis, published at ICLR'25. Samuele is a PhD student at the Dept. of Engineering, University of Oxford.

Angelo Huang
I co-supervised his Bachelor's thesis, published at EMNLP'24. Angelo is a Master's student in Computer Science at ETH.

Ping Zhu
I supervised his Master's thesis at Oxford. Ping is doing an MSc in advanced computer science at the University of Oxford.

Alberto Zurini
I co-supervised, with Alberto Cazzaniga, his Master's thesis. Alberto is a Master's student in Computer Science at the University of Udine.

Giovanni Monea
PhD student in Computer Science at Cornell University (2025-). *Lead the Future*

Simone Alghisi
PhD student in Information Engineering and Computer Science at the University of Trento. *Lead the Future*

Andrea Cerutti
Master's student in Computer Science and Engineering at the Polytechnic University of Milan. *Lead the Future*

Riccardo Inghilleri
Master's student in Computer Science at the Polytechnic University of Milan. *Lead the Future*

Orazio Torre
Bachelor's student in Computer Science at the University of Salerno. *Lead the Future*

Annalaura Pegoraro
Bachelor's student in math at the University of Padova and an exchange student at Berkeley (Summer 2024). *Lead the Future*

Mattea Busato
Bachelor's student at Bocconi University and an exchange student at the University of Toronto. *Lead the Future*

Christoph Weinhuber
PhD student at the University of Oxford, Dept. of Computer Science (2024-).

Saad Waheed
I supervised Saad for two terms for the Williams-Exeter exchange program (2024) - Saad is a bachelor's student at Williams College (US).

Alisa Kanganis
I supervised Alisa for two terms for the Williams-Exeter exchange program (2024) - Alisa is a bachelor's student at Williams College (US).

Simon Socolow
I supervised Simon for a term for the Williams-Exeter exchange program (2024) - Simon is a bachelor's student at Williams College (US).

Academic Service

Conference Reviewer: ICML, NeurIPS, ICLR, ACL, EMNLP.

Other Skills

Programming Languages: Python, C++. In the past, I used web languages (e.g., Javascript), and low-level languages (C and Assembly x86).

Libraries, Tools, Frameworks: PyTorch, CUDA (C++), Numpy, Eigen, Docker.